

Finite State, Forced Forgetting: Investigating the Memory Mechanism in Sub-quadratic Models

Anonymous ACL submission

Abstract

1 Introduction

2 Related Work

Memory in Transformers. Transformers implement in-context memory primarily through self-attention, where each token can retrieve information from earlier positions by computing content-dependent weights over key-value pairs (Vaswani et al., 2023). This mechanism is often viewed as a form of differentiable, content-addressable memory: rather than compressing the past into a single recurrent state, the model maintains access to a growing set of past representations, enabling flexible retrieval and composition over the entire context window. A large body of work has therefore focused on extending the effective context length or augmenting the model with additional memory beyond the fixed window. Notable approaches include recurrence across segments (Transformer-XL), which reuses previous hidden states to model dependencies beyond a fixed-length context (Dai et al., 2019), compressing older memories to preserve long-range information under a bounded compute budget (Rae et al., 2019), and retrieval-augmented language modeling, which explicitly consults an external datastore during generation (Borgeaud et al., 2022). These methods primarily target *memory availability* (how much information can be brought into the computation) and *access* (how that information can be retrieved when needed).

Attention-free sub-quadratic sequence models. A complementary line of work achieves sub-quadratic (often linear-time) inference by replacing attention with recurrent or convolutional updates that compress the past into a fixed-size state. State space models (SSMs) provide a particularly unified

view: they maintain an internal state updated online and produce outputs via a readout from that state. Foundational work on HiPPO introduced principled mechanisms for online compression of history via projections (Gu et al., 2020). Building on these ideas, structured parameterizations such as S4 made SSMs computationally practical and competitive on long-range tasks (Gu et al., 2022), and subsequent simplifications like DSS showed that diagonal state parameterizations can retain much of S4’s effectiveness (Gupta et al., 2022). Modern language-model-scale architectures (e.g., selective SSMs and gated linear attention variants) can often be interpreted as instances of this broader paradigm: the model stores information in a recurrent state and must *query* or *read* from it during generation. This raises a distinct memory question from the Transformer setting: when the past is compressed into a fixed-size state, the limiting factor may be less about what can be stored in principle and more about which stored information is practically retrievable given the model’s readout mechanism.

Entity binding as an in-context memory primitive. Entity-attribute binding tasks probe a core capability underlying factual recall and multi-entity reasoning: given a context describing several entities and their properties, the model must correctly bind each property to its corresponding entity at query time. Recent work has used controlled settings to study how binding is represented internally; for example, Feng and Steinhardt identify a “binding ID” mechanism and use causal interventions to argue that models attach binding-related vectors to entity and attribute positions (Feng and Steinhardt, 2024). Such findings motivate synthetic entity-binding evaluations as a targeted lens on in-context memory that avoids confounds from pretraining knowledge. Our entity-binding datasets are in this spirit, but we focus on sub-quadratic models whose memory is mediated by a recurrent state: this shifts

the emphasis from *Can tokens attend to the right binding to which bindings are present in the compressed state and how the model queries them*. In particular, by comparing what can be decoded from the state to what the model actually retrieves during generation, we align entity binding with a concrete hypothesis about the failure mode in sub-quadratic architectures: a potential mismatch between *stored* information and the model’s *query* mechanism.

Interpretability tools for studying memory mechanisms. Our work sits within mechanistic interpretability, which aims to localize and characterize internal computations responsible for model behavior. Probing methods train auxiliary predictors to decode information from hidden activations; recent ‘lens’ techniques decode intermediate representations into output distributions, providing a token-level view of how predictions evolve across depth (Belrose et al., 2025). Beyond correlational probes, causal interventions are commonly used to establish necessity and sufficiency of components. Meng et al. develop causal tracing interventions to localize computations supporting factual recall and show that specific mechanisms can be edited (ROME) (Meng et al., 2023). Activation patching has become a standard causal tool for circuit-level analysis, and recent work systematizes methodological choices and metrics for patching-based localization (Zhang and Nanda, 2024). Broader frameworks for circuit-based reasoning about Transformers provide a language for decomposing computations into interpretable parts (Elhage et al., 2021). We adopt these interpretability principles in the sub-quadratic setting by (i) designing probes that mirror the architecture’s native readout direction and (ii) using counterfactual patching on internal state representations to test whether probe-identified components have a causal effect on retrieval.

3 State Space Models

Many state-of-the-art sub-quadratic architectures, including Mamba [Add citation], RWKV [Add citation], GLA [Add citation], and Deltanet [Add citation], can be unified under the framework of *state space models* (SSMs); see section 3.1 for concrete examples. Despite their architectural differences, all of these models can be written in the following generic form.

The model maintains an internal state $S_t \in \mathbb{R}^{d \times d}$ that is updated at each time step based on the previ-

ous state and the current input $x_t \in \mathbb{R}^d$:

$$S_t = f(S_{t-1}) + g(x_t), \quad (1)$$

$$o_t = S_t q(x_t), \quad (2)$$

where f is a decay or transformation function applied to the previous state, $g(x_t)$ is some function of the current input (independent of other timesteps), and $o_t \in \mathbb{R}^d$ is the output. The output is produced by *querying* the state via right-multiplication with a query vector $q(x_t) \in \mathbb{R}^d$ derived from the current input.

Unlike transformers, whose attention mechanism incurs $O(L^2)$ computational cost with respect to sequence length L , SSM-based models enable inference in $O(L)$ time. This efficiency comes from compressing all past information into a fixed-size state $S_t \in \mathbb{R}^{d \times d}$, which is updated incrementally rather than recomputed from scratch.

This raises a fundamental question: given the finite capacity of the state, *how does the model store and retrieve in-context information?*

3.1 Examples

To illustrate the generality of this framework, we show how two popular architectures fit the SSM formulation.

Gated Linear Attention (GLA). In GLA, the state update and query take the form:

$$S_t = S_{t-1} \odot (1\alpha_t^\top) + v_t k_t^\top, \quad o_t = S_t q_t, \quad (3)$$

where $\odot$ denotes element-wise multiplication, and $\alpha_t, v_t, k_t, q_t \in \mathbb{R}^d$ are learned projections of the input.

Mamba-2. In Mamba-2, the formulation is:

$$S_t = \gamma_t S_{t-1} + v_t k_t^\top, \quad o_t = S_t q_t, \quad (4)$$

where $\gamma_t \in \mathbb{R}$ is a learned scalar decay factor.

Both architectures share the key property that the output is produced by querying the state from the right.

4 Methodology

4.1 Entity-Binding Datasets

We construct *in-context* entity-binding datasets to study how SSMs store and retrieve factual associations. Each dataset follows a generic structure: samples consist of N_e statements of the form “<entity> <relation> <property>”, where

Favorite Color ($N_e = 30, C = 10$) <i>Relation:</i> "X's favorite color is Y." Lady Gaga's favorite color is red. Beyoncé's favorite color is green. ...
Lives in City ($N_e = 30, C = 10$) <i>Relation:</i> "X lives in Y." The shark lives in France. The dog lives in Italy. ...

Figure 1: Examples from our entity-binding datasets. Each sample contains N_e entity-property pairs with one target entity whose property varies across samples.

each entity is associated with one of C possible property values.

These associations are provided entirely within the input prompt (not learned during pretraining) and vary randomly across samples. One designated *target entity* appears in all samples, but its associated property changes. The order of entities is randomized, and the target entity appears at a random position. This ensures that any successful retrieval must rely on in-context memory rather than memorized facts. See fig. 1 for examples.

4.2 Probing State Representations

Recall from eq. (2) that SSMs produce outputs by querying the state via right-multiplication: $o_t = S_t q(x_t)$. This suggests a natural way to probe whether entity-property bindings are stored in the state: we can train a probe that mimics this query mechanism.

For each internal state $S_t \in \mathbb{R}^{d \times d}$ (indexed by layer ℓ and head h), we extract the state at the *final token position* of the input sequence. We then train a probe with two learned components: (i) a **query component** $q_{\text{probe}} \in \mathbb{R}^d$ that queries the state from the right, producing $S_t q_{\text{probe}} \in \mathbb{R}^d$; and (ii) a **classifier component** $W_{\text{cls}} \in \mathbb{R}^{C \times d}$ that maps the queried vector to C class logits.

The full probe is:

$$\text{probe}(S_t) = W_{\text{cls}} S_t q_{\text{probe}} \in \mathbb{R}^C. \quad (5)$$

We call this the *natural probe* because it mirrors the model’s native state-querying mechanism: the state is multiplied by a vector *from the right*, just as in eq. (2).

To validate that this structure is important, we also train a *non-natural probe* that queries from the opposite direction:

$$\text{probe}_{\text{non-nat}}(S_t) = (q_{\text{probe}}^\top S_t) W_{\text{cls}} \in \mathbb{R}^C, \quad (6)$$

Given the following context, let’s answer the question below. Context: Lady Gaga’s favorite color is red. Beyoncé’s favorite color is green. ... Question: What is Lady Gaga’s favorite color? Answer: Lady Gaga’s favorite color is

Figure 2: Prompt template for evaluating model retrieval. The model must generate the correct property (e.g., “red”) as the next token.

where $q_{\text{probe}} \in \mathbb{R}^d$ is applied from the left, and $W_{\text{cls}} \in \mathbb{R}^{d \times C}$ projects to class logits. If memory is encoded consistently with the model’s query mechanism, the *natural probe* should significantly outperform the *non-natural probe*; we verify this in section 5.1.

4.3 Evaluating Model Retrieval

To compare probe accuracy with the model’s own retrieval ability, we construct a question-answering prompt that queries the model for the target entity’s property. fig. 2 shows an example for the Favorite Color dataset.

We measure top-1 accuracy: the model must generate the correct property label as its next token, where we only consider tokens corresponding to valid property values from the dataset (e.g., the $C = 10$ colors).

4.4 Causal Intervention

To establish that the heads identified by probing have a *causal* effect on model behavior, we perform counterfactual patching experiments using the prompt structure from section 4.3.

The procedure is as follows: (i) generate two prompts A and B that differ only in the target entity’s property (e.g., “Lady Gaga’s favorite color is red” vs. “Lady Gaga’s favorite color is blue”); (ii) run both prompts through the model, collecting states after processing the context; (iii) for prompt A , replace the states of a *subset* of heads with the corresponding states from prompt B ; and (iv) continue generation and measure whether the output switches from property A to property B .

We compare several head selection strategies: (i) **Baseline**: no heads patched, measuring the model’s normal accuracy; (ii) **Natural probe heads**: the top-ranked heads by *natural probe* accuracy; (iii) **Non-natural probe heads**: the top-ranked heads by *non-natural probe* accuracy; and (iv) **All heads**: re-

Table 1: Top 10 heads ranked by *natural probe* validation accuracy on the Favorite Color dataset.

Rank	Head	Val Acc.	Train Acc.
1	L15H11	96.3%	97.5%
2	L15H2	93.0%	95.4%
3	L15H13	91.4%	93.7%
4	L14H23	88.5%	91.4%
5	L15H27	79.0%	81.1%
6	L9H18	72.8%	76.6%
7	L15H1	65.3%	67.9%
8	L14H18	61.8%	65.9%
9	L15H8	61.1%	62.8%
10	L12H18	58.6%	63.3%

placing all head states, representing the maximum possible effect.

5 Results

We conduct all experiments using RWKV-7 1.5B, an SSM-based language model with 24 layers and 32 heads per layer (768 total heads), achieving state-of-the-art performance at its scale [Add citation].

5.1 Validating the Probing Approach

Before interpreting probe results, we first verify that the probe architecture is meaningful: it should only succeed when the target information is actually encoded in the state, rather than being powerful enough to extract arbitrary patterns.

5.1.1 The Probe is Not Overly Powerful

If the probe were too expressive, it could achieve high accuracy on any head regardless of whether that head stores relevant information. However, table 1 shows that the *natural probe* achieves high accuracy on only a small fraction of heads. The top head (L15H11) reaches 96.3%, but accuracy drops sharply: by rank 5 it falls below 80%, and by rank 10 below 60%. Most of the 768 heads yield near-chance performance.

This concentration of high-accuracy probes in a small subset of heads, rather than diffuse success across all 768 heads, indicates that the probe is identifying genuine “memory heads” rather than exploiting spurious correlations.

5.1.2 Querying from the Right Outperforms Querying from the Left

The *natural probe* (querying from the right) significantly outperforms the *non-natural probe* (querying from the left): the best *natural probe* reaches

Table 2: Counterfactual patching results (50 pairs). “Correct”: original property; “Switched”: counterfactual property; “Neither”: other output.

Method	Correct	Switched	Neither
Baseline	90%	0%	10%
<i>Natural</i> (top 0.5%)	14%	36%	50%
<i>Non-natural</i> (top 0.5%)	44%	4%	52%
All heads	0%	76%	24%

96.3%, while the best *non-natural probe* achieves only 70.8%.

This gap confirms that memory is encoded in a manner consistent with the model’s native query mechanism (eq. (2)). Probing from the right extracts information effectively; probing from the left does not. [TODO: add figure from poster]

5.2 Causal Effect of Memory Heads

To verify that probe-identified heads play a causal role in storing memory, we perform counterfactual patching experiments on 50 sample pairs from the Favorite Color dataset. table 2 shows the results.

Without patching (Baseline), the model outputs the correct property 90% of the time. When we patch the top 0.5% of heads selected by *natural probe* accuracy, the model switches to the counterfactual property 36% of the time. In contrast, patching the same number of heads selected by *non-natural probe* accuracy yields only a 4% switch rate. Patching all heads provides an upper bound, achieving 76% switch rate. [TODO: add experiment with random head selection as additional baseline]

5.3 The Bottleneck is Querying, Not Storage

A key observation emerges from comparing probe performance to model performance. While our probes extract the correct binding with up to 96.3% accuracy, the model’s own ability to answer questions about these bindings is substantially lower.

When we prompt the model using the template in fig. 2, its accuracy is far below probe accuracy. This gap implies that the information *is stored* in the state (the probe finds it) but the model struggles to *query* it effectively during generation.

This finding challenges the common assumption that SSMs’ limitations stem from insufficient memory capacity. Instead, our results suggest that the bottleneck lies in the model’s ability to formulate effective queries.

5.4 Universal Memory Heads

We analyze whether the same heads are important across different semantic domains by comparing top-performing heads on the Favorite Color dataset versus Lives in City.

We find substantial overlap: **7 of the top 10 heads are shared across both datasets**. This suggests the existence of *universal memory heads* that specialize in entity-property binding regardless of semantic content. We defer investigation of the mechanism underlying these heads to future work.

5.5 Semantic Collision Accelerates Forgetting

We investigate how context affects memory retention by measuring probe accuracy as a function of token distance from the information. We observe that **increased semantic similarity between the prompt content and the target subject accelerates the model’s forgetting rate**. When subsequent sentences discuss similar entities, probe accuracy degrades faster than when the intervening content is semantically distinct.

6 Future Work

Our findings open several directions for future research: (i) **Characterizing Universal Memory Heads**: what makes these heads special, and could understanding their mechanism inform architectural improvements? (ii) **Improving State Querying**: given that the bottleneck appears to be querying rather than storage, can we design better query mechanisms? (iii) **Scaling Analysis**: do these patterns hold in larger models?

7 Conclusion

We investigated the in-context memory mechanism of state space models by probing their internal states. Our key findings are: (i) **Natural probing works**: probes that query the state from the right significantly outperform probes that query from the left; (ii) **Memory is localized**: only a small subset of heads store entity-property bindings with high fidelity; (iii) **Probing identifies causal heads**: counterfactual patching confirms that probe-identified heads have genuine causal influence; (iv) **The bottleneck is querying**: probe accuracy far exceeds model accuracy, suggesting information is stored but poorly retrieved; (v) **Universal memory heads exist**: the same heads are important across different semantic domains; and (vi) **Semantic col-**

lision hurts: similar semantic content accelerates forgetting.

These results suggest that improving sub-quadratic models may require focusing on the query mechanism rather than storage capacity.

A Probing Training Details

Training configuration for the probing experiments: train/validation split of 70%/30%, maximum of 20,000 epochs, early stopping patience of 20 epochs, batch size of 1,000, and learning rate of 10^{-3} with Adam optimizer.

References

- Nora Belrose, Igor Ostrovsky, Lev McKinney, Zach Furman, Logan Smith, Danny Halawi, Stella Biderman, and Jacob Steinhardt. 2025. [Eliciting latent predictions from transformers with the tuned lens](#).
- Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George van den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, Diego de Las Casas, Aurelia Guy, Jacob Menick, Roman Ring, Tom Hennigan, Saffron Huang, Loren Maggiore, Chris Jones, Albin Cassirer, Andy Brock, Michela Paganini, Geoffrey Irving, Oriol Vinyals, Simon Osindero, Karen Simonyan, Jack W. Rae, Erich Elsen, and Laurent Sifre. 2022. [Improving language models by retrieving from trillions of tokens](#).
- Zihang Dai, Zhilin Yang, Yiming Yang, Jaime Carbonell, Quoc V. Le, and Ruslan Salakhutdinov. 2019. [Transformer-xl: Attentive language models beyond a fixed-length context](#).
- Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, et al. 2021. A mathematical framework for transformer circuits.
- Jiahai Feng and Jacob Steinhardt. 2024. [How do language models bind entities in context?](#)
- Albert Gu, Tri Dao, Stefano Ermon, Atri Rudra, and Christopher Re. 2020. [Hippo: Recurrent memory with optimal polynomial projections](#).
- Albert Gu, Karan Goel, and Christopher Ré. 2022. [Efficiently modeling long sequences with structured state spaces](#).
- Ankit Gupta, Albert Gu, and Jonathan Berant. 2022. [Diagonal state spaces are as effective as structured state spaces](#).
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. 2023. [Locating and editing factual associations in gpt](#).

- 415 Jack W. Rae, Anna Potapenko, Siddhant M. Jayaku-
416 mar, and Timothy P. Lillicrap. 2019. [Compressive](#)
417 [transformers for long-range sequence modelling](#).
- 418 Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob
419 Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz
420 Kaiser, and Illia Polosukhin. 2023. [Attention is all](#)
421 [you need](#).
- 422 Fred Zhang and Neel Nanda. 2024. [Towards best prac-](#)
423 [tices of activation patching in language models: Met-](#)
424 [rics and methods](#).